

Modern vs. traditional cryptography

Recap: traditional crypto

- An „art” rather than science
- Ad-hoc constructions
- Easy(ish) to break

Basic principles of modern crypto

- 1 formal and precise definitions
- 2 build on precisely stated assumptions
- 3 strictly proven security

Principle 1: precise definitions

Precise definitions

- design (What's the purpose / goal? Rather than „ex post facto”)
- usage (appropriate for goal?)
- analysis (comparison)
- intuition not enough

„Define” secure encryption

An encryption method is secure if ???

Principle 1: precise definitions

Definition of secure encryption

An encryption method is considered secure if

- 1 no attacker can recover the key from the ciphertext itself.
- 2 no attacker can recover the plaintext message from the ciphertext itself (even if only a small part of the message is missing)
- 3 no attacker can recover a single character of the message from the ciphertext (probabilities / order of magnitude of computation needed)
- 4 no attacker can learn anything important about the message knowing only the ciphertext (what counts as important?)
- 5 no attacker can recover any function (e.g. length, letter statistics, etc.) from the ciphertext

Principle 1: precise definitions

Definition of secure encryption

An encryption method is considered secure if

- 1 no attacker can recover the key from the ciphertext itself.
- 2 no attacker can recover the plaintext message from the ciphertext itself (even if only a small part of the message is missing)
- 3 no attacker can recover a single character of the message from the ciphertext (probabilities / order of magnitude of computation needed)
- 4 no attacker can learn anything important about the message knowing only the ciphertext (what counts as important?)
- 5 no attacker can recover any function (e.g. length, letter statistics, etc.) from the ciphertext

Principle 1: precise definitions

Definition of secure encryption

An encryption method is considered secure if

- 1 no attacker can recover the key from the ciphertext itself.
- 2 no attacker can recover the plaintext message from the ciphertext itself (even if only a small part of the message is missing)
- 3 no attacker can recover a single character of the message from the ciphertext (probabilities / order of magnitude of computation needed)
- 4 no attacker can learn anything important about the message knowing only the ciphertext (what counts as important?)
- 5 no attacker can recover any function (e.g. length, letter statistics, etc.) from the ciphertext

Principle 1: precise definitions

Definition of secure encryption

An encryption method is considered secure if

- 1 no attacker can recover the key from the ciphertext itself.
- 2 no attacker can recover the plaintext message from the ciphertext itself (even if only a small part of the message is missing)
- 3 no attacker can recover a single character of the message from the ciphertext (probabilities / order of magnitude of computation needed)
- 4 no attacker can learn anything important about the message knowing only the ciphertext (what counts as important?)
- 5 no attacker can recover any function (e.g. length, letter statistics, etc.) from the ciphertext

Principle 1: precise definitions

Definition of secure encryption

An encryption method is considered secure if

- 1 no attacker can recover the key from the ciphertext itself.
- 2 no attacker can recover the plaintext message from the ciphertext itself (even if only a small part of the message is missing)
- 3 no attacker can recover a single character of the message from the ciphertext (probabilities / order of magnitude of computation needed)
- 4 no attacker can learn anything important about the message knowing only the ciphertext (what counts as important?)
- 5 no attacker can recover any function (e.g. length, letter statistics, etc.) from the ciphertext

Principle 1: precise definitions

Definition of secure encryption

An encryption method is considered secure if

- ① no attacker can recover the key from the ciphertext itself.
- ② no attacker can recover the plaintext message from the ciphertext itself (even if only a small part of the message is missing)
- ③ no attacker can recover a single character of the message from the ciphertext (probabilities / order of magnitude of computation needed)
- ④ no attacker can learn anything important about the message knowing only the ciphertext (what counts as important?)
- ⑤ no attacker can recover any function (e.g. length, letter statistics, etc.) from the ciphertext

Principle 1: precise definitions

How to make the definition formal?

- What does „break a cipher” / „recover the message” mean?
- What does „no attacker” mean (what powers do they possess)?

Example

A cryptographic protocol meant for a specific purpose is secure if no attacker with the specified (computational) power can perform a specified form of attack.

Principle 1: precise definitions

Math vs. practice

- hardware-based attacks
- human factors

Good definition of security/secretcy should

- support the intuitive view
- be supported by examples
- be backed up by ongoing analysis over time

Principle 2: precise assumptions

Two variants

- unconditional secrecy
- computational secrecy

Why?

- validation of assumptions
- comparison of methods
- facilitate security proofs

Principle 3: formal proofs of secrecy

Why?

- established difficulty vs. naive intuition
- works \nRightarrow unbreakable
- risks of poor cryptosystem or poor software product

Proof of secrecy by reduction

Protocol X is considered secret (by a certain definition) if assumption Y is correct.

Perfect secrecy

Informálisan

What do we need to specify a scheme?

- three algorithms: Gen, Enc, Dec
- message space \mathcal{M}

Parameters

- *key space*: \mathcal{K} set of possible keys ($k \in \mathcal{K}$)
- *message space*: \mathcal{M} set of possible messages ($m \in \mathcal{M}$)
- *ciphertext space*: \mathcal{C} set of possible ciphertexts
- Usually finite (esp. keyspace - „large” but finite)

Perfectly secret scheme

Parameters

- *Probability distributions over $\mathcal{K}, \mathcal{M}, \mathcal{C}$*
- $k \in \mathcal{K} : Pr(K = k)$ denotes the probability that key k is chosen.

e.g.: \mathcal{K} : bit sequences of length 128,
 $k \in_R \mathcal{K} \Rightarrow Pr(K = k) = 1/2^{128}$

- Similarly for \mathcal{M}, \mathcal{C}

e.g.: $|\mathcal{M}| = 2, Pr(\text{Attack tomorrow}) = 0.7, Pr(\text{No attack}) = 0.3$

- Distribution over \mathcal{K} and \mathcal{M} independent and arbitrary
- Distribution over \mathcal{C} determined by the other two.
- *conditional probability: $Pr(A | B)$: Probability of A , provided that we know B is true.*

Perfectly secret scheme

Parameters

- *Probability distributions over $\mathcal{K}, \mathcal{M}, \mathcal{C}$*
- $k \in \mathcal{K} : Pr(K = k)$ denotes the probability that key k is chosen.

e.g.: \mathcal{K} : bit sequences of length 128,
 $k \in_R \mathcal{K} \Rightarrow Pr(K = k) = 1/2^{128}$

- Similarly for \mathcal{M}, \mathcal{C}

e.g.: $|\mathcal{M}| = 2, Pr(\text{Attack tomorrow}) = 0.7, Pr(\text{No attack}) = 0.3$

- Distribution over \mathcal{K} and \mathcal{M} independent and arbitrary
- Distribution over \mathcal{C} determined by the other two.
- *conditional probability*: $Pr(A | B)$: Probability of A , provided that we know B is true.

Perfectly secret scheme

Definition

A scheme is a triple $\Pi = (Gen, Enc, Dec)$ where :

- *Gen* is key generation, a probabilistic algorithm that returns a key $k \in_R \mathcal{K}$ (maybe using an input called the security parameter)
- *Enc* is encryption, a probabilistic algorithm that returns a ciphertext $c \in \mathcal{C}$ on inputs $k \in \mathcal{K}$ and $m \in \mathcal{M}$, i.e.
 $c := Enc_k(m)$.
- *Dec* is decryption a deterministic algorithm that returns a plaintext upon inputs k and $c \in \mathcal{C}$: the return value is $Dec_k(c)$ in \mathcal{M} .

Secrecy definitions

Intuition

- we know the distribution of messages
- knowing the ciphertext, no information about the message should be learnt
- attacker's computational power: infinite

Definition

A scheme over \mathcal{M} is perfectly secret if for any distribution over \mathcal{M} and $\forall m \in \mathcal{M}, \forall c \in \mathcal{C}$:

$$Pr(M = m) = Pr(M = m | C = c).$$

Secrecy definitions

Equivalent formulation

- We cannot distinguish the ciphertexts corresponding to two different messages.

Lemma (Perfect indistinguishability)

A scheme provides perfect secrecy if for any distribution over \mathcal{M} , and $\forall m_1, m_2 \in \mathcal{M}, \forall c \in \mathcal{C}$:

$$Pr(C = c | M = m_1) = Pr(C = c | M = m_2).$$

Secrecy definitions

Equivalent formulation

- indistinguishability game: two players: adversary and tester

Indistinguishability experiment with eavesdropper $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}}$

- 1 Adversary \mathcal{A} issues messages m_0, m_1 with $|m_0| = |m_1|$
- 2 Tester randomly chooses key k and bit $b \in_R \{0, 1\} : c = \text{Enc}_k(m_b)$. Send c to \mathcal{A} .
- 3 \mathcal{A} answers by outputting $b' \in \{0, 1\}$
- 4 $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}} = 1$ if $b = b'$, otherwise 0. (Wins the game if guesses correctly.)

Secrecy definitions

Indistinguishability experiment with eavesdropper $\text{PrivK}_{\mathcal{A},\Pi}^{\text{eav}}$

- 1 Adversary \mathcal{A} issues messages m_0, m_1 with $|m_0| = |m_1|$
- 2 Tester randomly chooses key k and bit $b \in_R \{0, 1\} : c = \text{Enc}_k(m_b)$. Send c to \mathcal{A} .
- 3 \mathcal{A} answers by outputting $b' \in \{0, 1\}$
- 4 $\text{PrivK}_{\mathcal{A},\Pi}^{\text{eav}} = 1$ if $b = b'$, otherwise 0. (Wins the game if guesses correctly.)

Definition

A scheme Π is perfectly secret over \mathcal{M} if $\forall \mathcal{A}$:

$$\Pr(\text{PrivK}_{\mathcal{A},\Pi}^{\text{eav}} = 1) = \frac{1}{2}.$$

Lemma

These definitions are equivalent.

One-time pad

One-time pad (OTP)

Initialize $\mathcal{K} = \mathcal{M} = \{0, 1\}^n$

Gen let $k \in_R \{0, 1\}^n$ uniformly random

Enc for $k \in \{0, 1\}^n$ and $m \in \{0, 1\}^n$, let
 $c = Enc_k(m) = k \oplus m$.

Dec for k and $c \in \{0, 1\}^n$ let $Dec_k(c) = c \oplus k$.

Theorem

One-time pad is perfectly secret.

Drawbacks of perfect secrecy

One-time pad properties

- $|k| = |m| \Rightarrow$ too long keys / short messages only
- "one-time" (really, never reuse!)
- these are not unique to OTP, but inherent to perfect secrecy

Theorem

Let Π be a perfectly secret scheme over \mathcal{M} and let \mathcal{K} be the key space determined by Gen . Then $|\mathcal{K}| \geq |\mathcal{M}|$.